WEEK 1 · COURSE 1 · 2 HOURS · ALL PERSONNEL

AI Fluency Fundamentals

The six 201-level skills that separate sustained adopters from the 80% who quit.

EXPERT-DRIVEN DEVELOPMENT · MCD-MONTEREY

SPEAKER NOTES

Hook: Welcome them in. Wait for the room to settle. Read the title slowly: "AI Fluency Fundamentals." Then say: "This is the first of six weeks. Everyone in the program takes this one — Marines, sailors, civilians, contractors. We're going to spend two hours together."

Emphasize: "We are not teaching you a tool today. We are teaching you how to manage one." Pause on that line. It is the thesis of the entire course.

Bridge: "Before we get into the six skills, I want to tell you why most people fail at this — and why you won't."

BEFORE WE START

How this works.

- **Two hours, six modules, one ten-minute break** at the midpoint.
- **Two hands-on activities** — the Red Pen Review and a workflow mapping exercise. Have a notepad open.
- **Stop me with questions.** If something is unclear here, it will be unclear in your section meeting tomorrow.
- **You'll need an AI tool.** GenAI.mil if you have CAC; ChatGPT, Gemini, or CamoGPT also work. Open it now.
- **This is unclassified.** No real names, no PII, no operational details in your prompts today.

SPEAKER NOTES

Cue: Pause and watch the room open their AI tool. Don't move on until you see most laptops light up.

Emphasize: "Stop me with questions" — mean it. Repeat it. Quiet rooms are confused rooms.

Engagement: Quick poll — "Show of hands, who has used GenAI.mil before? ChatGPT? Gemini? Never used any of them?" Calibrate your examples for the room.

Bridge: "OK. Here's the agenda."

AGENDA

Six modules. Two hours.

0:00 – 0:15	**1 · Why 80% Quit**	Frontier Recognition
0:15 – 0:40	**2 · The Six Skills That Actually Matter**	All six — overview
0:40 – 0:55	**3 · The Delegation Equation**	Task Decomposition, Context Assembly
0:55 – 1:05	**Break**	10 minutes
1:05 – 1:30	**4 · The Trust Problem — Red Pen Review**	Quality Judgment, Frontier Recognition
1:30 – 1:45	**5 · Centaur, Cyborg, or Neither**	Workflow Integration
1:45 – 2:00	**6 · Frontier Mapping & Your Assignment**	Frontier Recognition, all six

SPEAKER NOTES

Hook: "Here's where we're going. Six modules, two hands-on activities in the middle, ten-minute break."

Emphasize: Module 4 is the Red Pen Review. "If you only remember one thing from today, that's where you'll get it." Don't preview the trick — just flag it as the highlight.

Bridge: "Module 1 is fifteen minutes. Let's go."

01

MODULE ONE

Why 80% Quit (and how you won't)

15 MIN

201 SKILL: FRONTIER RECOGNITION

SPEAKER NOTES

Cue: Section dividers exist so the room knows we're shifting gears. Pause for two beats.

Hook: "Module One. Fifteen minutes. We're going to talk about why most people fail at AI."

Bridge: "Microsoft has data on this — hundreds of thousands of employees."

80%

MICROSOFT WORK TREND INDEX

Three weeks of excitement. Then a crater.

Microsoft tracked AI tool adoption across 300,000+ employees. Excitement peaked at three weeks. Then most people quietly stopped.

SPEAKER NOTES

Hook: "Microsoft has the largest dataset on AI adoption in the world. Three hundred thousand employees, real workplaces, tracked over months."

Emphasize: "Eighty percent quit. Not because the tool was bad. Because nobody taught them how to use it." Land on "nobody taught them."

Engagement: "How many of you have tried an AI tool, given up, and never gone back? Be honest." Hands will go up.

Bridge: "Here's exactly what that failure looks like — you'll recognise it."

THE FAILURE LOOP

It always looks the same.

- You type "help me with this report." You get something generic.
- You try again. You get something confident and wrong.
- You try a third time. You decide it's faster to do it yourself.
- You never come back.

SPEAKER NOTES

Cue: Read each line slowly. Pause between bullets. Let the recognition land.

Emphasize: "Confident and wrong" is the killer. Most people quit at step two because they realise they can't trust it. We will fix that today.

Engagement: "Anyone here had this exact arc? Show of hands. OK — you are normal."

Bridge: "Here's what happens when people get the right training."

25 MIN

UK GOVERNMENT DIGITAL SERVICES

Per worker. Per day. With training.

20,000 government workers, 12 departments, three months. After training, 80% refused to give the tools back — and 9 of 12 departments kept their licenses.

SPEAKER NOTES

Hook: "Here's the flip side. The UK government is the closest analogue to us — bureaucratic, hierarchical, security-conscious."

Emphasize: 25 minutes a day is two weeks a year. Per worker. The math is enormous. But notice — "with training." Not "with the tool."

Bridge: "So what kind of training? Not what you'd think."

“

You are not learning to use a tool. You are learning to manage one.

THE THESIS OF THIS COURSE

SPEAKER NOTES

Cue: Stop. Read it once, slowly. Pause. Read it once more.

Emphasize: "If you take nothing else from today, take that line. Manage, not use." This is the spine of every other slide in this deck.

Bridge: "Here's why that framing matters. You already know how to manage."

THE MANAGEMENT FRAMING

Would you hand a 100-page RFP to a brand new Marine and say "handle this"?

- No. You'd **break it into pieces**.
- You'd **tell them which parts to tackle first**.
- You'd **explain what good looks like**.
- You'd **review their work** and give feedback.

That's how you should work with AI.

SPEAKER NOTES

Hook: Read the headline as a real question, not rhetorically. Wait. Get a couple of "no" responses from the room.

Emphasize: The four bullets are the four skills you already use every day with your Marines. AI is just another team member who needs the same management.

Engagement: "What's the difference between a great NCO and a bad one? It's not the tools. It's what's on this slide."

Bridge: "Quick check, then we move into the six skills."

KNOWLEDGE CHECK — MODULE 1

What separates the 20% who stick with AI from the 80% who quit?

> **Not** better prompts. **Not** a better tool. **Not** being more technical.
>
> They **manage** AI — like a capable but inexperienced team member. They give it context, break work into pieces, evaluate the output, and iterate. The same skills you already use to lead Marines.

SPEAKER NOTES

Cue: Show the question. Wait. Don't show the answer for at least 20 seconds.

Engagement: Cold-call one student: "What do you think?" Then a second. Then click for the reveal.

Emphasize: The framing word is **manage**. Hammer it.

Bridge: "OK. Module 2. Twenty-five minutes. The six skills."

02

MODULE TWO

The Six Skills That Actually Matter

25 MIN ALL SIX 201 SKILLS — OVERVIEW

SPEAKER NOTES

Hook: "Module Two. Twenty-five minutes. We're going to walk through all six skills. For each one I'll show you the 101 behaviour and the 201 behaviour side by side."

Emphasize: "The gap between 101 and 201 is where the value lives. Pay attention to the gap."

Bridge: "Here are the six."

THE SIX 201 SKILLS

Six skills. None of them are prompt formulas.

01

Context Assembly

Give AI the right information to produce useful output.

02

Quality Judgment

Know whether AI output is correct, complete, and appropriate.

03

Task Decomposition

Break large tasks into AI-appropriate subtasks.

04

Iterative Refinement

Improve AI output through structured feedback loops.

05

Workflow Integration

Embed AI into recurring processes, not one-off experiments.

06

Frontier Recognition

Know what AI can and cannot do — and when that boundary shifts.

SPEAKER NOTES

Cue: Don't read all six aloud — let them scan. Give them 15 seconds.

Emphasize: "Notice what's not on this list: nothing about prompt syntax, nothing about which model is best, nothing about temperature settings. Those don't matter at this level."

Bridge: "Skill one. Context Assembly."

SKILL 1 · CONTEXT ASSEMBLY

Same task. Same AI. Different prompts.

101 — VAGUE

```
Write me a counseling statement.
```

201 — EXPERT

```
Write a counseling statement for a Lance Corporal who was
late to formation twice this month. Tone: corrective but not
adversarial. The Marine is otherwise a solid performer. Use
the format from NAVMC 10274. Page 11 entry, not an adverse
6105.
```


SPEAKER NOTES

Hook: "Read both. The expert prompt has five things the vague one doesn't — rank, situation, tone, format, document type."

Emphasize: "AI is extremely sensitive to context. The difference between mediocre and useful is almost always the information **you** provided."

Engagement: "What do you think the AI does with the vague prompt? Right — it makes up a Marine, makes up a situation, gives you something generic. Garbage in, garbage out applies."

Bridge: "Skill two. Quality Judgment. This one matters more than the rest combined."

Same prompt. Two outputs. Which gets approved?

OUTPUT A

```
SSgt Smith demonstrated exceptional leadership and
unparalleled dedication to duty. His outstanding performance
and tireless work ethic resulted in tremendous success
across all areas of responsibility...
```

OUTPUT B

```
SSgt Smith led a four-Marine maintenance team responsible
for 87 principal end items valued at $2.1M, maintaining 98%
operational readiness over 18 months with zero
accountability discrepancies across three command
inspections...
```


SPEAKER NOTES

Hook: "Same AI. Same prompt — an end-of-tour NAM for a Staff Sergeant. Both came back instantly. Which one survives the awards board?"

Engagement: Show of hands — A or B? Cold-call one defender from each side **before** you pick.

Emphasize: "B has facts — 87 items, $2.1 million, 98 percent. A has filler — **exceptional**, **tremendous**, **multi-million**. AI defaults to language that sounds impressive over language that is accurate. Your job is rejecting **sounds good** and demanding **is accurate**."

Bridge: "Skill three. Decomposition."

SKILL 3 · TASK DECOMPOSITION

Don't hand it the whole task. Hand it the next step.

101 — ONE GIANT ASK

```
Write me a training plan for the platoon.
```

201 — SEQUENCED SUBTASKS

```
First, list the required annual training events for FY26.
Then map them against the calendar, avoiding known
conflicts. Then identify which events need range time and
put those first.
```


SPEAKER NOTES

Hook: "Same task. The first prompt asks for everything at once. The second breaks it into three sequential steps with clear inputs and outputs."

Emphasize: "Smaller chunks are more likely to succeed. This is exactly how you'd give a task to a junior Marine. You don't say **handle the whole training plan**. You say **start with the FY26 events**."

Bridge: "Skill four. Iteration."

SKILL 4 · ITERATIVE REFINEMENT

First draft is 70%. Three passes get you to 95%.

- **Pass 1.** "The tone is too formal. Make it sound like a Gunny talking to his Marines."
- **Pass 2.** "Add the duty officer info: Capt Rodriguez, (831) 555-0147. Add a 250-mile liberty radius. Cut the third paragraph — it repeats the first."
- **Pass 3.** "Tighten to one page. Add: no swimming at Point Lobos due to rip current advisory. Remind them about the Monday 0630 formation."

Three specific prompts. Not "make it better." **Specific**.

SPEAKER NOTES

Hook: "This is a real example. Weekend safety brief. AI's first draft was a generic blob — 70% there. Three refinement passes turned it into something a Gunny would actually read at formation."

Emphasize: "Notice what every refinement prompt has in common — they're **specific**. Never **make it better**. Always **change this exact thing to this exact other thing**." This is the same feedback you'd give a junior Marine on a draft.

Engagement: "How many of you accept the first draft AI gives you and walk away? OK — stop doing that. The first draft is the starting line, not the finish line."

Bridge: "Skill five."

Stop using it sometimes. Start using it on the same task every week.

101 — SIDE ACTIVITY

> "I'll try the AI thing for this one report later."

201 — RECURRING STEP IN THE WORKFLOW

> "Every Friday the duty NCO drafts the weekend safety brief in GenAI.mil from the current weather, liberty boundaries, and recent incidents. That's how we do it."

SPEAKER NOTES

Hook: "The 101 use is a side experiment. The 201 use is a standing step in a recurring workflow. Same tool. Completely different outcome."

Emphasize: "Microsoft's research says it takes about 11 weeks to build the AI habit. If you only use it once a month, you'll never get past the crater of disappointment we talked about in Module 1."

Engagement: "Pick one weekly task in your section right now. That's your candidate for next week's homework."

Bridge: "Last skill. Frontier Recognition. This is the one that keeps you out of trouble."

Know what AI can do for you. Know what it can't. Know that the line moves.

AI HANDLES WELL

Drafting correspondence, summarising long documents, formatting, naming conventions, brainstorming, restructuring text, translating between styles.

AI HANDLES POORLY

Interpreting specific MCOs, calculating TIS/TIG accurately, anything requiring institutional knowledge that isn't in its training data, anything where the cost of being wrong is high.

SPEAKER NOTES

Hook: "These two columns are not the same length. The boundary between them is what researchers call the jagged frontier."

Emphasize: "The boundary is **not intuitive**. Tasks that look simple sometimes fail. Tasks that look hard sometimes succeed. You have to discover where the frontier is for **your** work — nobody can hand you a complete map."

Engagement: "When you discover a task AI handles poorly, that knowledge is as valuable as discovering what it handles well. **Share it** with your section."

Bridge: "Here's the cost of getting Frontier Recognition wrong."

−19 PTS

BCG × HARVARD, 2023

Untrained AI use can make you *worse* than no AI at all.

When 758 BCG consultants used AI on tasks *outside* the frontier, they were 19 percentage points less likely to produce correct work than the control group with no AI.

SPEAKER NOTES

Hook: "This is the number that should scare you a little. AI doesn't just fail to help on the wrong tasks — it actively makes you **worse**."

Emphasize: "These are smart people. Top-tier consultants. They got 19 points **worse** because they trusted AI on tasks where it should not be trusted. That's the cost of skipping skill six."

Bridge: "OK. Quick check, then we're into Module 3 — the delegation equation."

An AI draft of an awards write-up has dates, names, dollar figures, and reads beautifully. Sign it?

> **No.** Not until **Quality Judgment** kicks in. AI presents fabricated facts with the same confidence as real ones. Verify the dollar figures against the property book. Verify the dates against the award period. Verify the reference numbers exist.
>
> Reads beautifully ≠ is accurate. Your name goes on the bottom.

SPEAKER NOTES

Cue: Read the question. Pause for at least 15 seconds. Cold-call.

Emphasize: "Your name. Goes on the bottom. AI is not signing this — you are."

Bridge: "Module 3. Fifteen minutes. The Delegation Equation. This is how you decide whether to bother with AI on any given task."

03

MODULE THREE

The Delegation Equation

15 MIN TASK DECOMPOSITION · CONTEXT ASSEMBLY

SPEAKER NOTES

Hook: "Module Three. Fifteen minutes. We need a way to decide which tasks are worth handing to AI and which aren't. Ethan Mollick at Wharton built one. It's three questions."

Bridge: "Here it is."

THREE QUESTIONS, EVERY TASK

Before you delegate, ask:

- **1. Human baseline.** How long would this take me to do myself?
- **2. Probability of success.** How likely is AI to produce acceptable output?
- **3. AI process time.** How long does it take me to request, wait, and *evaluate*?

Delegate when (1) is large, (2) is high, and (3) is much smaller than (1).

SPEAKER NOTES

Hook: "Three numbers. That's the whole equation."

Emphasize: "Number three is the trap. People forget evaluation time. If AI saves you four hours of writing but costs you four hours of fact-checking, you broke even — and you trusted AI for nothing."

Bridge: "Three quick examples. Watch how the math changes."

Verdict: Delegate.

- **Human baseline:** 4 hours from scratch.
- **AI drafts in:** minutes.
- **Review time:** 45 minutes for accuracy and unit-specific details.
- **Probability AI gets the structure right:** ~70%.
- **Net:** roughly three hours saved. **Delegate the structure, fix the specifics.**

SPEAKER NOTES

Hook: "Five-paragraph order. Most of you have written one. Let's run the math."

Emphasize: "AI is great at the structure — SITUATION, MISSION, EXECUTION, ADMIN, COMMAND. It's bad at the specifics. So delegate structure, keep specifics."

Bridge: "Now a counter-example."

Verdict: Don't delegate.

- **Human baseline:** 5 minutes with the right references.
- **AI process:** may or may not be right, hard to tell.
- **Review time:** as long as just doing it yourself.
- **Net:** zero saving, real risk of error. **Do it yourself.**

Short tasks with a high cost of being wrong are usually not worth delegating.

SPEAKER NOTES

Hook: "Same person, different task. The math says don't delegate. AI isn't the answer to everything."

Emphasize: "If checking takes as long as doing, the AI bought you nothing. And the cost of getting a Marine's promotion math wrong is real."

Bridge: "One more — and this one's a hybrid."

Verdict: Delegate the structure. Fill in the specifics yourself.

- **Human baseline:** 6 hours.
- **AI produces:** a solid first-draft structure in minutes.
- **Review time:** 30 minutes to align with your unit's calendar.
- **Net:** ~5 hours saved on structure. Specifics still come from you.

SPEAKER NOTES

Hook: "Most of your real work looks like this third example. Partial delegation."

Emphasize: "**The structure is delegable. The specifics aren't.** That's a useful pattern. Hold onto it."

Bridge: "There's one more thing about this equation that matters — and it's the reason your seniority makes you better at AI, not worse."

"

AI doesn't replace expertise. It rewards it.

THE EXPERTISE MULTIPLIER — MOLLICK, 2026

SPEAKER NOTES

Hook: "The more expertise you have, the better you are at all three numbers in the delegation equation. You give better instructions. You catch errors faster. You know which tasks are worth delegating."

Emphasize: "If you're a senior NCO worried that AI replaces what you bring — the research says the opposite. Your judgment is the multiplier."

Bridge: "Quick check, then we take a ten-minute break."

KNOWLEDGE CHECK — MODULE 3

A task takes you 10 minutes. AI does it in 30 seconds — but you spend 12 minutes verifying. Delegate?

> **No.** The third number in the equation — AI process time — includes evaluation. If checking takes longer than doing, the AI cost you time, not saved you time.
>
> The fast tasks where AI *looks* efficient are often the ones where the math doesn't actually work. **Always count evaluation time.**

SPEAKER NOTES

Cue: Read the question. Pause. Cold-call. **Then** reveal.

Emphasize: "Speed is a trap. Always count evaluation time."

Bridge: "Ten-minute break. We're back at [state the wall-clock time] for the Red Pen Review."

HALFWAY POINT

Break.

Ten minutes. Stretch, refill, check messages. We're back at [TIME] for the Red Pen Review.

MODULE 4 STARTS AT THE TOP OF THE HOUR

SPEAKER NOTES

Cue: Write the actual return time on the board / in chat: "We resume at HH:MM." Don't trust the room to track ten minutes themselves.

Logistics: While they're on break, pull up your three Red Pen Review documents and verify the share link works. If the room is on Teams, post the documents in the chat now so latecomers have them.

Bridge: "Coming back — this is the most important 25 minutes of the course."

04

MODULE FOUR

The Trust Problem — Quality Judgment

25 MIN

QUALITY JUDGMENT · FRONTIER RECOGNITION

SPEAKER NOTES

Hook: "Welcome back. Module Four. Twenty-five minutes. This module pays for the entire course."

Bridge: "Here's the problem we have to solve."

AI is confidently right and confidently wrong — in the same paragraph.

THE FUNDAMENTAL QUALITY JUDGMENT PROBLEM

SPEAKER NOTES

Hook: "The problem isn't that AI lies. The problem is that AI doesn't **know** when it's lying. Real facts and made-up facts get the same confident tone."

Emphasize: "AI doesn't flag uncertainty. It doesn't say **I'm not sure about this part**. That's **your** job."

Bridge: "We're going to practice that job right now. Three documents."

EXERCISE 10 MIN INDIVIDUAL · 5 MIN PAIR · 10 MIN DEBRIEF

Red Pen Review.

YOUR JOB

01 Read each of the three AI-generated documents in the Teams chat.

02 Mark every claim you would **not sign** — facts, references, numbers, procedures, anything.

03 Note *why* you flagged each one.

USE THIS CHECKLIST

- References & citations — do they exist?
- Facts & statistics — verifiable?
- Internal consistency — do dates / names / units match?
- Procedural accuracy — is this how it really works?
- Substance vs style — real content or polished filler?
- Would you **sign it?**

SPEAKER NOTES

This is a hands-on slide — stop clicking.

Cue: Drop the three documents into the Teams chat now. Set a 10-minute timer visible to the room.

Engagement: While they work, walk the room (or scan reactions in Teams). After the individual work, ask them to compare with one neighbour for 5 minutes — not to argue, just to see what the other person flagged that they missed.

Emphasize: "Finding 3–4 issues per doc is doing well. Finding 1 or zero just taught you why this exercise exists."

Bridge: "Time. Let's debrief — doc by doc."

THE THREE DOCUMENTS

Realistic AI output. Not trick questions.

- **Doc 1 · Award narrative.** A NAM recommendation for a Corporal — reads polished, looks reasonable.
- **Doc 2 · SOP excerpt.** Check-in / check-out procedures with reference numbers and step lists.
- **Doc 3 · Training write-up.** An after-action style summary with statistics and dates.

These are not "gotcha" documents with one planted error. They're realistic first drafts — the kind of output AI actually produces, with the kinds of problems it actually creates.

SPEAKER NOTES

Cue: Use this slide as a transition into the debrief sequence. Don't read it — orient them on what's coming.

Emphasize: "These are not trick documents. The errors here are organic — inflated language, fabricated references, invented statistics, procedural gaps. Exactly what AI gives you in real life."

Bridge: "Doc 1. What did you find?"

DEBRIEF · DOCUMENT 1

Award narrative for Cpl Hernandez.

WHAT'S WRONG

→ "47% reduction" and "156 man-hours per quarter" are unsupported — AI invented them.

→ Reference to "MARADMIN 045/26" — verify it exists; cited MARADMINs are a common AI fabrication.

→ "Highest mark in the regiment" — bold claim, no source. Awards boards will ask.

→ Filler vocabulary: *unparalleled*, *tireless*, *selfless*. Sound impressive, prove nothing.

WHAT'S RIGHT

→ Format follows NAM narrative conventions.

→ Period of service and unit are correctly placed.

→ Closing reflects-credit-upon paragraph is correct boilerplate.

SPEAKER NOTES

Cue: Before revealing, ask the room: "Who flagged the 47% number? Who flagged the MARADMIN reference?" Get a count.

Emphasize: "The made-up MARADMIN is the dangerous one. AI invents reference numbers in confident, plausible formats. If you don't look them up, you sign a fabrication."

Bridge: "Doc 2 is more subtle."

DEBRIEF · DOCUMENT 2

SOP excerpt — check-in / check-out.

WHAT'S WRONG

- → Reference (a) "MCO 1000.6A" — verify the order and revision letter exist.
- → "CMC Form 4790/142" — AI loves to invent plausible form numbers. Look it up.
- → Times and locations ("Bldg 1284, Rm 203", "every Monday at 0800") — specific enough to look real, but ungrounded in any source you gave the AI.
- → Procedural gap: no mention of the medical, dental, or finance check-in stops most units actually require.

WHAT'S RIGHT

- → Document structure and numbering follow standard SOP conventions.
- → Inclusion of an effective date and references section is correct.
- → The 72-hour reporting window is a reasonable default.

SPEAKER NOTES

Cue: Ask: "Who looked up the form number? Who flagged the procedural gaps?" Most rooms catch one but not both.

Emphasize: "Procedural gaps are the most dangerous category. Made-up references can be caught with a search. Missing steps require **you** to know how the process really works. There is no substitute for domain knowledge."

Bridge: "Doc 3."

Training event after-action.

WHAT'S WRONG

- → Statistics presented to one decimal place ("98.7%") suggesting precision the source data can't support.
- → Inconsistent attribution — the document credits one billet for work another billet performed.
- → Claims that are positive but unverifiable: "highest in the regiment", "first time in five years".
- → Recommendations section parrots the observations — no real analysis underneath.

WHAT'S RIGHT

- → Five-paragraph after-action structure is correctly applied.
- → Tone is appropriate — descriptive, not adversarial.
- → Useful as a starting structure if you replace the invented specifics with real ones.

SPEAKER NOTES

Cue: "What did you flag in Doc 3? Who caught the wrong attribution?" Pull from at least two students.

Emphasize: "AI loves false precision. **98.7 percent** sounds verifiable. **About 99 percent** sounds estimated. AI defaults to the first because it sounds more authoritative — even when there's no data behind it."

Bridge: "There's a pattern across all three. Here's the lesson."

Match the depth of review to the cost of being wrong.

- **High-stakes documents** (legal, personnel, financial) — line-by-line verification of every fact, reference, and number. Centaur mode. Human verifies everything.
- **Medium-stakes documents** (correspondence, reports, briefings) — spot-check key claims, verify references, check tone. Mixed mode.
- **Low-stakes documents** (first drafts, brainstorming, formatting) — quick review for obvious errors. Cyborg mode.

SPEAKER NOTES

Hook: "Not every document deserves the Red Pen Review. But a personnel package **always** does. Match your review depth to the stakes."

Emphasize: "Centaur and cyborg are coming up in Module 5 — remember those words. They're the two patterns for working with AI."

Bridge: "Quick check. Then we move on."

An AI document includes a confident reference to "MCO 5215.1J". You don't recognize it. Action?

Look it up. Always. AI invents Marine Corps order numbers in confident, plausible formats — correct prefix, correct format, completely fabricated content.

If a reference can't be verified in MCO Library / EPME / your reference library, it's fabricated. Strike it. **Verifiable references go in. Made-up ones don't.**

SPEAKER NOTES

Cue: Question. Pause. Cold-call.

Emphasize: "If you don't recognise a reference, look it up. Always. The number of times this single habit has saved someone an embarrassing signature is enormous."

Bridge: "Module 5. Centaur, Cyborg, or Neither. Fifteen minutes."

05

MODULE FIVE

Centaur, Cyborg, or Neither

15 MIN

WORKFLOW INTEGRATION

SPEAKER NOTES

Hook: "Module Five. Fifteen minutes. There are two productive patterns for working with AI — both work, both are research-backed. The 201 skill is knowing which one fits which task."

Bridge: "Here they are, side by side."

Pick the pattern that fits the task. Don't pick one and apply it to everything.

CENTAUR

Clear division of labour. Human does Phase 1, hands off to AI for Phase 2, human reviews in Phase 3. Distinct responsibilities, clear handoff points, verification at every boundary.

Best for: high-stakes work that needs accountability. Legal review, admin separations, security packages, fitrep drafts.

CYBORG

Continuous integration. Constant back-and-forth between human and AI. The boundary is fluid — sometimes you write, sometimes AI does, sometimes you edit each other.

Best for: creative and iterative work. Building Power Apps, writing SOPs, drafting training materials, exploratory data analysis.

SPEAKER NOTES

Hook: "BCG and Harvard ran 758 consultants through this. Both patterns outperformed unstructured AI use. The mistake was picking one pattern and applying it to everything."

Emphasize: "If accountability matters — centaur. If creativity matters — cyborg. Same person, both modes, different tasks."

Bridge: "Let's look at each one a little closer."

Phase. Handoff.
Phase. Handoff.

- **Human designs.** You decide what's needed, what good looks like, and the rules.
- **AI executes.** AI drafts, formats, summarizes — whatever the design calls for.
- **Human verifies.** Every fact, every reference, every number — line by line.
- **AI refines.** You hand back specific corrections; AI applies them.
- **Human accepts.** Final review before your name goes on it.

Verification at every boundary. This is the pattern for anything you sign.

SPEAKER NOTES

Hook: "Centaur is structured. Five distinct phases, with a verification step between each AI pass."

Emphasize: "Use this for fitreps, separations, awards packages — anything where the cost of an error is high."

Bridge: "Cyborg looks completely different."

CYBORG, IN PRACTICE

Continuous back-and-forth. Both hands on the wheel.

- Start with a rough idea: *"I need to see who's overdue on annual training."*
- Build incrementally — data model, input, dashboard, reports.
- At each step, evaluate and redirect: *"That chart shows it by person; I need it by date."*
- Sometimes you type. Sometimes AI generates. The boundary is fluid.

Use this for creative work, prototyping, and exploration — not for documents you sign.

SPEAKER NOTES

Hook: "Cyborg is fluid. There are no phases. You and AI are basically pair-programming, even if you're not coding."

Emphasize: "Cyborg works for the messy front end of any project — building a tool, drafting an SOP, exploring data. Once you have something to **sign**, you switch back to centaur for the verification."

Bridge: "Now let's get specific. Here's the activity."

EXERCISE 7 MIN INDIVIDUAL · 3 MIN SHARE-OUT

Map one of your workflows.

ON YOUR NOTEPAD

01 Pick **one recurring task** from your real job.

02 Break it into 3–5 subtasks.

03 Mark each subtask: **Human only** / **AI could help** / **AI should do this**.

04 Decide: is this **centaur** or **cyborg**?

05 Estimate time saved if AI handled the appropriate subtasks.

THEN

- Two volunteers will share their map.
- The room reacts: did they put a subtask in the wrong column?
- Common pattern: **30–50%** of recurring work has AI-appropriate subtasks people have never tried delegating.

SPEAKER NOTES

This is a hands-on slide — stop clicking.

Cue: 7-minute timer. While they work, post the prompt in chat: **"Recurring task · 3-5 subtasks · Human / Could help / Should do this · Centaur or Cyborg · Time saved."**

Engagement: Pick two volunteers in advance if you can. If nobody volunteers, name two students directly. Don't let silence kill the exercise.

Emphasize at debrief: "The 25 minutes a day from the UK study didn't come from one big win. It came from dozens of small workflow integrations. **Compounding**."

Bridge: "Quick check. Then last module."

You're prototyping a Power App. Centaur or cyborg?

Cyborg. Prototyping is creative, iterative, exploratory — the boundary between you and AI is fluid. You're not signing anything; you're discovering what the right tool looks like.

Once the prototype works and you're documenting it for production — switch back to centaur. Verify everything before your name goes on the user guide.

SPEAKER NOTES

Cue: Question. Pause. Cold-call. Reveal.

Emphasize: "Same project, both modes. Cyborg for the messy front end. Centaur for the polished output. Pattern-switching is the 201 move."

Bridge: "Last module. Frontier mapping. Fifteen minutes."

06

MODULE SIX

Frontier Mapping & Your Assignment

15 MIN FRONTIER RECOGNITION · ALL SIX

SPEAKER NOTES

Hook: "Module Six. Last fifteen minutes. We're going to start a frontier map for your unit and I'm going to give you an assignment."

Bridge: "Reminder — the frontier is the line between what AI can and can't do. It's jagged, and it moves."

The line is uneven. The line moves. Map it for your work.

AI HANDLES WELL TODAY

Drafting correspondence · summarizing long documents · restructuring text · brainstorming · formatting · explaining concepts · translating between styles.

AI HANDLES POORLY TODAY

Specific MCO interpretation · TIS/TIG calculations · current personnel data · anything classified · anything where institutional knowledge isn't in its training set.

SPEAKER NOTES

Hook: "Two columns. They will both change next quarter. The one on the right gets shorter every model release. Your map has to be a living document."

Emphasize: "There's a third column we should add — **moving frontier**. Tasks where capability is improving and you should re-test periodically. Set a recurring calendar reminder. **Re-test what AI couldn't do six months ago**."

Bridge: "Here's the in-class activity to start your section's map."

EXERCISE 5 MIN INDIVIDUAL · 5 MIN SHARE-OUT

Seed your section's frontier map.

ON YOUR NOTEPAD

01 Write three columns: **Handles well**, **Handles poorly**, **Moving frontier**.

02 Add at least **two examples** per column from your real work.

03 Star the failure cases that surprised you.

THEN

- Three volunteers share one example from each column.
- The room adds to the master list in the Teams chat.
- That list goes home with you — it's your section's starter map.

SPEAKER NOTES

This is a hands-on slide — stop clicking.

Cue: 5-minute timer. Open a shared note in Teams chat titled "[Unit] Frontier Map — Seed". Paste columns into it.

Engagement: As students share, you transcribe into the shared note. Three volunteers, one example per column each.

Emphasize: "Failure cases — the starred ones — are the most valuable thing you can give your section. **Sharing what AI got wrong saves your peers an hour of confusion.**"

Bridge: "Time. Here's your homework."

YOUR ASSIGNMENT — BEFORE NEXT WEEK

Three things. One week.

- **1.** Pick one recurring task from your workflow map (Module 5).
- **2.** Try AI on it this week — centaur or cyborg, your call.
- **3.** Bring your results — especially your **failure cases** — to share with your section.

If you're continuing to Builder Orientation: bring a problem you want to **build a tool to solve**.

SPEAKER NOTES

Cue: Read the three items. Don't skip the failure-case bullet — it's the most important.

Emphasize: "Failure cases are not embarrassments. They're data. The Microsoft research says it takes 11 weeks to build the AI habit. You're on Day 1. Failure is part of the curve."

Bridge: "Let's recap the whole course in one slide."

COURSE RECAP

What you take with you.

6

Skills. Context, Quality, Decomposition, Iteration, Workflow, Frontier. None are prompt formulas.

3

Numbers in the equation. Human baseline. Probability of success. AI process time — including evaluation.

2

Patterns. Centaur for what you sign. Cyborg for what you explore. Switch as the task changes.

1

Habit. Always look up references. Always count evaluation time. Always share failure cases.

11

Weeks. The time it takes to build the AI habit. You are on Day 1. Stay on the curve.

80%

Quit. They didn't have what you have now. You will be in the other 20.

SPEAKER NOTES

Cue: Don't read the cells. Let the room scan. Pause for 20 seconds.

Emphasize: "The 80 in the bottom-right corner — that was the opening number. You are not in that 80 anymore. You have the framework that puts you in the 20."

Bridge: "Here's where you go from here."

WHERE YOU GO FROM HERE

If you want to use AI better — you're done. If you want to *build* — Course 2 is next.

COURSE 2 · BUILDER ORIENTATION

2 hours. Aspiring builders. From user to builder. Live builds. Debugging with AI. Decomposition exercises on a real problem you bring.

Bring: a problem in your workflow you want to solve with a tool.

FOR EVERYONE ELSE

Keep applying the six skills this week. Re-take this course on the EDD site any time. Add to your section's frontier map. Recommend the course to one person who hasn't taken it.

SPEAKER NOTES

Hook: "Two paths from here. Both are good."

Emphasize: "You don't have to take Course 2. AI Fluency Fundamentals is a complete course on its own. But if you've ever thought **I bet I could build a tool for that** — sign up."

Bridge: "Last slide. Then we open the floor."

INSTRUCTOR ONLY **AFTER CLASS — BEFORE YOU FORGET**

What to do once the room clears.

TODAY (WITHIN 1 HOUR)

01 Send the Teams chat transcript — especially the seeded frontier map — to every attendee.

02 Note which knowledge-check questions caused the longest pause; flag for next delivery.

03 Capture any failure cases the room shared — they're examples for the next cohort.

THIS WEEK

01 One-week follow-up message: ask attendees what they tried and what failed.

02 Forward Builder Orientation sign-ups to the program coordinator.

03 Update the EDD frontier-map page with new examples from the cohort.

04 Schedule the next session before momentum decays. **11 weeks to habit, remember.**

SPEAKER NOTES

Audience: This slide is for **you**, not the room. Skip past it during a live delivery, or hide it by pressing End early. It exists so the deck doubles as your post-class checklist.

Why it matters: The crater of disappointment isn't only an attendee problem — it's an instructor problem too. If you don't follow up within a week, attendance was for nothing.

Bridge: Press next once for the closing / Q&A slide.

THANK YOU

Questions?

You now have the framework that separates the 20% who stick from the 80% who quit. The tools will keep changing. The six skills won't.

EXPERT-DRIVEN DEVELOPMENT · MCD-MONTEREY · UNCLASSIFIED

SPEAKER NOTES

Cue: Open the floor. Don't fill silence — let the room ask.

If silence: Prompt with: "What's the first task you're going to try AI on this week?" That's an easier question than "Any questions?" and it surfaces the real concerns.

Close: Thank them by name if it's a small room. Remind them about Course 2 and where to find materials on the EDD site. Wrap on time — respect of the clock is part of the demo.