Expert-Driven Development

WEEK 3 OF 6 · COURSE 3

Platform Training

Build three complete tools on Power Platform. Practice centaur and cyborg work patterns under real time pressure.

4 hours · Builders **Prerequisite:** Builder Orientation · M365 access

SPEAKER NOTES

OPEN Welcome the room. Confirm everyone has a laptop, M365 credentials, and is signed in to GenAI.mil or CamoGPT. Today is the longest session in the program — four hours, three live builds. Set expectations: low-text slides, lots of platform time, and break placement after the first build.

Tone: this is the working session. Builder Orientation was the on-ramp. Today we ship.

LOGISTICS Tell the room how to follow along, where the bathroom is, and that they can press **N** on a presenter laptop to pull up these notes if they're ever co-presenting.

WHERE WE ARE

Weeks 1 & 2

You've already done the hard part.

WEEK 1 · AI FLUENCY

- The six 201 skills
- The jagged frontier
- Quality judgment over prompting
- The delegation equation

WEEK 2 · BUILDER ORIENTATION

- From user to builder
- Your first prototype
- Task decomposition in practice
- The problem you brought to today

SPEAKER NOTES

RECAP Two weeks ago you learned the management skills. Last week you turned them into a build. Today you do it three times in a row, on the platform you'll use back at your unit.

Ask quickly: "Show of hands — who deployed something from Builder Orientation already?" If many hands go up, congratulate them and tell them today builds on that. If few, say "today is where it sticks."

PIVOT "Everything we do today is the same skill set, applied at production speed. Centaur, cyborg, frontier recognition — we'll exercise all of it."

TODAY'S MISSION

By 4 hours from now

Three working tools. One updated frontier map.

- **Build #1 — Centaur:** a request routing workflow you'd trust in production.
- **Build #2 — Cyborg:** a training tracker discovered as you build it.
- **Build #3 — Your problem:** the tool you brought from Week 2, in the mode you choose.
- **One frontier-map update:** at least one new entry in the unit's shared map of where AI helps and where it hurts.

SPEAKER NOTES

STAKES Make it concrete. By the time we wrap, you'll have shipped three things and added a row to the unit's frontier map. None of these is a toy build.

Reinforce the through-line: Build #1 forces structure. Build #2 forces fluidity. Build #3 forces you to choose the right tool. The frontier map is what you take back to your section.

WATCH THE ROOM If anyone looks anxious about pace, name it: "We've got 60 minutes per build, you'll have prompt scaffolds for the first two, and I'll circulate the whole time."

THE 4-HOUR MAP

Pace yourself

Three builds. Two breaks. One map update.

0:00	**Module 1 — Setup & Review**	15 min · All six skills
0:15	**Build #1 — Request Routing (Centaur)**	60 min · Context, Quality
1:15	*Break + Failure Sharing*	15 min · Frontier
1:30	**Build #2 — Training Tracker (Cyborg)**	60 min · Iteration, Workflow
2:30	*Break*	15 min
2:45	**Build #3 — Your Problem (Choose Mode)**	60 min · Decomposition, all six
3:45	**Frontier Map Update & Wrap**	15 min · Frontier

SPEAKER NOTES

PACE Walk the timeline once, top to bottom. Don't dwell. The point is that they see the rhythm: build, debrief, break, build, debrief, break, build, debrief, map.

Call out the break placement explicitly: ***"Our long break is at 1:15 — that's also when we do failure sharing. We share failures during the break because that conversation is more honest with coffee in hand."***

BUFFER The course budget is ~3h20m of instruction with a 40-minute buffer inside the 4-hour block. Don't say "we have buffer" — say "we'll move." Use the buffer privately if a build runs over.

MODULE 1 OF 6

Setup & Review.

15 minutes

Goal: zero access surprises in Build #1

SPEAKER NOTES

SECTION Brief stop. The goal of this module is brutal: confirm everybody can reach SharePoint, Power Apps, Power Automate, and a CAC-enabled AI tool, and confirm Integrate works in their tenant. If we discover an access problem during Build #1, we lose 20 minutes recovering. We discover it now.

Resolve before Build #1

Six checks. None optional.

PLATFORMS

- SharePoint Online — can create a list
- Power Apps — `make.gov.powerapps.us`
- Power Automate — `high.flow.microsoft.us`
- GenAI.mil or CamoGPT — can start a conversation

CAPABILITIES

- **Integrate button** visible in any SharePoint list
- **Approvals connector** appears in Power Automate

If Integrate is missing, build canvas apps from blank instead. If Approvals is missing, skip Build #1 Steps 5–6.

SPEAKER NOTES

RUN THE CHECK Have everyone open all four URLs right now and physically click into the Integrate menu on a test SharePoint list. Then have them click **+ New flow** in Power Automate and search "Approvals."

Hands up if anything fails. Triage in this order: (1) can't reach a URL → IT support, (2) Integrate missing → use the canvas-app fallback prompt, (3) Approvals connector missing → flag now, we'll skip the optional approval step in Build #1.

TIME CHECK Spend at most 5 minutes here. If someone is fully blocked on access, pair them with a working partner for the first build and resolve in the background.

Two patterns. Pick the one that fits the build.

CENTAUR

Clear phases. Human designs and verifies. AI executes between checkpoints.

- **Best for:** high-stakes, accuracy-critical workflows
- **Risk:** slower, but reliable
- **Today:** Build #1 — request routing

CYBORG

Continuous back-and-forth. Boundary between human and AI work stays fluid.

- **Best for:** exploratory, requirements-discovered-as-you-build work
- **Risk:** faster, but demands constant attention
- **Today:** Build #2 — training tracker

SPEAKER NOTES

REFRESH One-minute review. We're not relearning this — we're calibrating the room before we apply it twice.

Ask 2–3 students: "Last week's assignment — what mode did you use, and did it fit the problem?" Use one short answer to land the point that mode-choice is itself a 201 skill.

ANCHOR Today: Build #1 forces centaur. Build #2 forces cyborg. Build #3 you choose. By the end you'll have lived in both.

Pick before you prompt

Use the highest-trust tool your data allows.

CAC-ENABLED (PREFERRED)

- **GenAI.mil** — CUI authorized; Agent Designer available
- **CamoGPT** — CUI authorized; paste the agent primer at session start

COMMERCIAL (UNCLASSIFIED ONLY)

- ChatGPT · Gemini — OK for non-CUI exercises and demo prompts
- PII stays anonymized regardless of tool, unless a PIA authorizes otherwise

SPEAKER NOTES

DECIDE Have the room call out which tool they're using today. If most are on GenAI.mil with Agent Designer, great. If they're on CamoGPT, remind them to paste the quick-start primer at the top of every conversation — see the Agent Setup Guide on the EDD site.

Reinforce: CUI is OK on GenAI.mil and CamoGPT. PII still gets anonymized. Commercial tools are for unclassified work only — Build #3 with real unit data must use the IL5 path.

TIP Tell them the Power Apps trick: open Tree View, right-click any control, Copy. Paste that into the AI conversation when they need precise modifications. We'll use this on every build.

BUILD #1 OF 3 · MODULE 2

Centaur Mode. Request Routing.

60 minutes

Five phases. Four checkpoints.

SPEAKER NOTES

SECTION Hard pivot into the first build. Energy up. From here until 1:15, we're heads-down on a real workflow.

Frame the centaur pattern out loud one more time: ***"You design. AI builds. You verify. AI refines. You accept. At every boundary, you check before you cross."***

Before you touch the platform.

THE PROBLEM

Route a request to the right approver, get a decision back, notify the requester.

1st Bn, 99th Marines needs a request workflow. Submit → auto-route by dollar amount → approver acts → requester notified. SharePoint list, Power App form, Power Automate flow.

MODE

Centaur — rules and approvals matter

High-stakes accuracy. Wrong approver = real consequences. Phase boundaries are non-negotiable.

SUCCESS LOOKS LIKE

A $400, $1,200, and $3,000 test request each route correctly, and the requester gets the right notification.

Bonus: approver timeout escalates after 48 hours.

SPEAKER NOTES

FRAME Read all three boxes aloud. The point of this slide is that everyone in the room can repeat back: **what we're building, why centaur fits, and what done looks like.**

If a student can't repeat those three things, they're not ready to start. Pause and re-frame.

ANTI-PATTERN TO CALL OUT "Don't open Power Apps yet. We design first. The single biggest failure mode in centaur is starting to build before the rules are written down."

BUILD #1 · PHASE 1

10 minutes · No AI yet

Whiteboard first. AI never.

Define every business rule, every data flow, every edge case. On paper or a whiteboard. The quality of your design decides the quality of the final product.

- **Business rules** — Who approves what, at what threshold?
- **Data flow** — Form → list → notification → action → notification
- **Notifications** — Requester, approver, admin — what does each one need?
- **Edge cases** — Approver on leave, request withdrawn, duplicate submission

SPEAKER NOTES

CUE **Stop clicking. Pick up a marker.** Have students whiteboard or paper-write the four sections above for ten minutes — alone or in pairs. No AI, no platform. Just rules.

Walk the room. Look for two failure modes: (1) someone has skipped straight to the platform, and (2) someone has written "the approver approves" and called it a rule. Push back: ***"What is the threshold? What's the alternate? What happens at 11:59 PM on day two?"***

SAMPLE ANSWERS If someone is stuck, share the canonical answers from the instructor page: under \$500 → Section Leader, \$500–\$2000 → Dept Head, over \$2000 → CO; 48-hour timeout escalates; alternate approver if primary is on leave. But make them write it themselves.

BRIDGE "Once your rules are on the board, you're ready to feed them to AI. Next slide is your switch cue."

SWITCH TO POWER PLATFORM — NOW

Open SharePoint. Build the list. Then Integrate it.

We leave the deck for the next ~40 minutes. I'll walk you through the canonical prompt sequence on the projector. Open your AI tool side-by-side with Power Apps.

SharePoint · Leave Requests list | Power Apps · Integrate → Create app | Power Automate · routing flow | Approvals · if connector available

---

SPEAKER NOTES

SWITCH **Stop advancing slides. Switch your screen-share to Power Platform.** From here you teach live. Walk through the four canonical prompts (SharePoint design → form customization → routing flow → approval action) with the room building alongside you.

Order of operations on screen:

1. Show the SharePoint list creation, paste their whiteboard rules into the AI for the CSV template, import. 2. Click **Integrate → Power Apps → Create an app**. Paste the EditForm customization prompt. 3. Click **Integrate → Power Automate → Create a flow**. Paste the routing prompt. 4. If Approvals is available, replace the "Send email" with "Start and wait for an approval" and set timeout to P2D.

MID-BUILD CHECK AT STEP 3 Pause at the end of the routing flow and ask: ***"What parts of the AI's suggestions did you have to verify or correct so far?"*** Take 2–3 answers. Don't move on without at least one.

BUILD #1 · DEBRIEF

5 minutes

What centaur cost you. What it bought you.

What it cost

- Ten minutes on a whiteboard before any code
- A verification checkpoint at every phase boundary
- Slower than “just ask AI to build it”

What it bought

- A workflow you'd actually deploy to your section
- Errors caught at the boundary, not in production
- A design document, in your handwriting, that survives the build

SPEAKER NOTES

DEBRIEF Three short questions, one minute each. Take quick verbal answers from the room — no slides, no prompts, no “let me think.”

1. ***"At which phase boundary did you catch the most AI errors?"***
2. ***"What did the AI try to do that the whiteboard caught?"***
3. ***"Where would you have been if you'd skipped Phase 1?"***

LAND THE POINT "This is what centaur buys you. Slow at the front, sturdy at the back. Now we go to break — and during the break, we share what AI got wrong."

MODULE 3 · DURING THE BREAK

Failure Sharing.

15 minutes total

Quiet, open, attributable

SPEAKER NOTES

SETUP Frame this carefully before the break. ***"This is one of the most important fifteen minutes in the entire program. We're going to share what AI got wrong in Build #1, and how we caught it. No grades. No judgment. Every failure case you share saves someone else an hour of confusion."***

Lead by sharing your own AI failure first. Be specific: a wrong DAX function, a hallucinated SharePoint column name, a routing rule it inverted. Set the tone that experienced builders fail too.

Move to the quiet layout next slide. Encourage people to grab coffee, then come back to share.

OPEN DISCUSSION · 10 MINUTES

What did AI get wrong? How did you catch it? What will you check next time?

Share one specific failure case. Skip the framing — we already know the build. Tell us what you saw and how you fixed it.

— The exact AI output that was wrong

— The phase, step, or test where you noticed

— What you told the AI to fix it

— One line you're adding to your verification checklist

SPEAKER NOTES

RUN THE DISCUSSION Stand back. Let the room talk. Aim for at least four people sharing — start with anyone who's already volunteered, then specifically invite a quieter participant if needed.

Capture failure cases on a whiteboard or shared doc as people share. Group them by type: ***"AI invented a column"***, ***"AI swapped two rules"***, ***"AI's syntax was off"***. The grouping is the start of the unit's frontier map — you'll come back to this list in Module 6.

PSYCHOLOGICAL SAFETY If the room is quiet, don't fill the silence. Re-prompt with: ***"What's the dumbest thing AI did to you in the last hour? Bonus points for embarrassing."*** A laugh usually opens it up.

END At the 10-minute mark, name two patterns you heard, then hand off to the break splash.

BREAK

5 minutes

Stretch, refill, then we open Build #2 at the cyborg pace.

SPEAKER NOTES

BREAK Five minutes of true downtime after failure sharing. Project the time on the wall clock. Reset your own screen-share to Power Platform with a fresh, empty conversation in your AI tool — Build #2 starts fast.

BUILD #2 OF 3 · MODULE 4

Cyborg Mode.
Training Tracker.

60 minutes

No phases. Continuous conversation.

SPEAKER NOTES

SECTION Energy back up. We're shifting modes deliberately. Where Build #1 made you wait at every checkpoint, Build #2 wants you in continuous conversation. You'll feel the difference in the first ten minutes.

BUILD #2 · PRE-BUILD FRAMING

Cyborg

Discover the requirements as you build.

THE PROBLEM

"I need to see which Marines have completed which training events."

Start with the rough need. Build a SharePoint data model, an input form, a Power BI dashboard with conditional formatting, and a calculated days-until-due column. Restructure the data when the dashboard tells you to.

MODE

Cyborg — the right shape isn't obvious yet

You'll change your mind about the data model at least once. That's the point.

SUCCESS LOOKS LIKE

A dashboard that color-codes status across five training events, and you can explain why you reshaped the data mid-build.

Bonus: a weekly S-3 readiness email driven by a scheduled flow.

SPEAKER NOTES

FRAME Read all three boxes. Emphasize that the data restructure is intentional. The first model — one row per training event — won't give you a good dashboard. You'll discover that, throw it out, and rebuild as one row per Marine. That's cyborg working at full speed.

CALIBRATION "In Build #1 you weren't allowed to touch the platform until the rules were written. In Build #2 you start typing prompts in 90 seconds. Notice how different that feels."

SWITCH TO POWER PLATFORM — NOW

Open a fresh AI conversation. Start typing.

No whiteboard. No outline. Tell the AI the rough need, evaluate what comes back, and redirect. The boundary between you and the AI stays fluid for the next ~45 minutes.

SharePoint · Training Status list | Iterate the data model | Power BI · conditional formatting | Calculated columns · days until due

SPEAKER NOTES

SWITCH **Stop advancing slides. Switch your screen-share to Power Platform.** Live-walk the six iterations from the instructor page:

1. SharePoint list (one row per event — the wrong model). 2. Restructure to one row per Marine (the cyborg pivot). 3. Power BI with conditional formatting (this is where the AI gives wrong instructions — let it). 4. Fix the conditional formatting after students hit the wall. 5. Calculated days-until-due column. 6. Optional: scheduled S-3 readiness email.

DELIBERATE FAILURE The AI's first conditional-formatting answer is wrong. **Do not correct it.** Let students fail. When hands go up, ask: ***"What did the AI tell you to do? What result did you get? What would you ask it to clarify?"*** This is frontier recognition in real time.

MID-BUILD CHECK AT ITERATION 3 Pause and ask: ***"How is this build process different from Build #1? What feels***

BACK TO DECK · BUILD #2 DEBRIEF

BUILD #2 · DEBRIEF

5 minutes

What cyborg cost you. What it bought you.

What it cost

- You threw away your first data model
- You had to spot a wrong AI answer in real time
- You couldn't walk away from the conversation

What it bought

- You shipped a working dashboard in 60 minutes
- The shape of the data emerged from the build, not from a meeting
- You practiced fail-and-pivot at speed

SPEAKER NOTES

DEBRIEF Three quick questions:

1. ***"At what point did you realize the first data structure wasn't going to work?"***
2. ***"How did Build #2 feel different from Build #1?"***
3. ***"If you had to do Build #2 over, what's the one thing you'd do differently?"***

LAND THE POINT "Cyborg is faster, but it's not safer. It works because you're paying attention every second. The minute you stop watching, it goes off the rails."

Hand off to the second break.

BREAK

15 minutes

Bigger break this time. Eat. Walk. Come back ready to build your own thing.

SPEAKER NOTES

BREAK Long break. Reset clocks, refill water, and quietly check in with anyone who fell behind during Build #2 — pair them with a working partner for Build #3 if needed.

Use this time to look around the room: is anyone visibly struggling? Anyone visibly bored? Calibrate Build #3 difficulty in your head before you start.

BUILD #3 OF 3 · MODULE 5

Your Problem.
Choose Your Mode.

60 minutes

Real data · real workflow · real users

SPEAKER NOTES

SECTION Tone shift. The first two builds were yours to run, but the rules were mine. This one is yours all the way down. Real problem. Real users. Your call on the mode.

Reinforce the data-handling reminder: real unit data is OK on GenAI.mil and CamoGPT. PII gets anonymized. Commercial tools are off-limits for CUI.

Your call

State three things out loud. Then start.

DECLARE TO THE ROOM

One sentence. One mode. Two frontier risks.

Before you open the platform, say what you're building, which mode you're using and why, and where you expect AI to struggle. If you can't name two specific frontier risks, you're not ready — revise.

WHAT YOU'RE BUILDING

Name the tool type and the user. ***"A vehicle dispatch tracker for the motor pool chief."***

MODE & WHY

Tie the mode to the task. ***"Centaur, because the approval chain has strict business rules."***

SPEAKER NOTES

RUN DECLARATIONS Go around the room. Each person gets 30 seconds. **You assess in under a minute:** is the scope realistic for 60 minutes, is the mode choice defensible, are the frontier risks specific?

Push back hard on vagueness. ***"Something to help with vehicles"*** → revise. ***"AI might struggle"*** with no specifics → revise. The declaration is itself a 201 skill — the act of stating it sharpens the build.

CUE SUCCESS CRITERIA For each person: "What does done look like at 3:45?" If they can't say it, they don't have it.

SWITCH TO POWER PLATFORM — NOW

Build the thing you came here to build.

I'll circulate. When you hit a block, ask me what you'd tell the AI — not how to fix it. Teaching you to debug with AI is more valuable than fixing it for you.

- CUI on GenAI.mil or CamoGPT only
- Anonymize PII unless PIA-authorized
- Choose centaur or cyborg deliberately
- Stop at 3:40 — debrief in deck

SPEAKER NOTES

SWITCH **Stop advancing slides. Switch to Power Platform and circulate the room.** No live build on screen this time — the room is heads-down on their own work.

HOW TO COACH When a student hits a wall, resist the urge to give the answer. Instead ask: ***"What would you tell the AI about what's going wrong?"*** Then pause. Let them rephrase the prompt. The skill is the rephrasing.

WATCH FOR Students who haven't typed in 5 minutes (stuck — go help). Students who haven't paused to test in 20 minutes (going too fast — make them verify). Students who declared centaur but jumped straight to building (call it out).

TIME DISCIPLINE At 3:30, give a 10-minute warning. At 3:40, return everyone to the deck for debrief.

BACK TO DECK · BUILD #3 DEBRIEF

BUILD #3 · DEBRIEF

5 minutes

Did your mode choice survive contact with the build?

Look back

- What did your mode choice actually buy you?
- Did you switch modes mid-build? If so, when and why?
- Where on the frontier did your AI fail?

Look forward

- What's the smallest next step to make this deployable?
- Who's the first user you'd put it in front of?
- What's the verification protocol before they touch it?

SPEAKER NOTES

DEBRIEF Two-question rotation around the room. Each person answers **one question from the left and one from the right.** Keep it tight — 30 seconds per answer. We need to be at the frontier-map module by 3:50.

CAPTURE As people speak, jot down the failure cases on the same whiteboard you started during failure sharing. We're going to fold them into the unit frontier map next.

MODULE 6 · WRAP

Frontier Map. Then we close.

15 minutes

One real entry per person

SPEAKER NOTES

SECTION Final stretch. The frontier map is the thing today produced that survives past today. Three builds end up on the floor when you leave; the map goes back to your section.

Use this shape for your own map

The unit frontier map.

TASK TYPE	AI HANDLES WELL	AI HANDLES POORLY	VERIFICATION NEEDED
Form & list scaffolding	CSV templates, EditForm formulas, default values	Person field claims format, choice field defaults	Submit test items at every threshold
Routing & approval flows	Conditional branches, email body templates	Inverted thresholds, missing timeout settings	Trace every branch with named test cases
Power BI visuals	Card visuals, table layout, basic DAX	Conditional formatting steps, chart-type fit	Re-do with one column at a time, sanity-check chart against data
Unit-specific knowledge	Generic policies, standard formats	Specific MCO numbers, TIS/TIG, local SOPs	Cross-check against the actual reference document
Your row from today →	*add what worked*	*add what failed*	*add the check that catches it*

SPEAKER NOTES

WALK THE MAP Read the four columns out loud once, then walk the four pre-filled rows. Tell the room: ***"This is what an EDD frontier map looks like. Four columns, one row per task type. The last row is what you're about to add."***

Keep this slide up while students start drafting their own row on the next slide. They should be looking at this layout while they write.

MODULE 6 · WORKTIME

Add at least one row

Write the row.

Pull from your three builds today. Pick the failure case that surprised you most. Write the row using the four columns from the previous slide.

7 MINUTES

SPEAKER NOTES

WORKTIME Set a literal seven-minute timer. Project the previous slide on a side screen if you can — students need to see the column headings while they write.

Walk the room. The most common failure mode here is generic entries. ***"AI sometimes makes mistakes"*** is not a row. ***"AI inverted the dollar-amount routing thresholds in Power Automate; caught at the $1,200 test; verification = trace every threshold with a named test case"*** is a row.

At the end, ask three volunteers to read their row out loud. Capture the rest in the shared map. Tell them: ***"GySgt Swickard pushes the consolidated map MCCES-wide. Your row helps everyone."***

WRAP

Four hours later

What you can now do.

- Build a Power Platform tool in **centaur mode** with verification at every phase boundary.
- Build a Power Platform tool in **cyborg mode**, including a deliberate data-model pivot.
- Choose between modes for your own problem, defend the choice, and ship in 60 minutes.
- Recognize a wrong AI answer fast enough to course-correct in the same prompt.
- Add a real, specific row to your unit's frontier map.

SPEAKER NOTES

LAND THE WRAP Read the list once, slowly. Then say: ***"This is what you couldn't do at 0900. You can do it now."***

Don't oversell. The room knows what they shipped. Acknowledge that the third build was harder than the first two and that some people are still mid-flight on it. That's expected.

BEFORE WEEK 4

Required to attend Advanced Workshop

Four things between now and Week 4.

- **Finish Build #3** to deployable state.
- **Run it through the EDD SOP QA process** — not skipping the verification checkpoints.
- **Document three specific failure cases** you encountered while finishing.
- **Identify one capability surprise** — somewhere AI performed better or worse than you expected.

SPEAKER NOTES

ASSIGNMENT This is the entry ticket to Week 4. Advanced Workshop assumes you've shipped at least one tool through the SOP and have failure cases to bring. Without those, the Week 4 exercises don't land.

LOGISTICS Tell them where to log failure cases (the shared frontier-map doc on the EDD site) and who to contact if they get blocked finishing Build #3 between sessions.

WEEK 4 OF 6 · COURSE 4

Advanced Workshop.

Four hours for experienced builders. You bring a deployed tool and the failure cases you collected this week. We map the frontier for your domain and build verification protocols you can hand to someone else.

- Frontier mapping for your specific domain
- Multi-component system build with named QA gates
- Verification protocols: write them, hand them off
- Group debugging on real, in-flight problems
- Teaching others — the 201 multiplier
- Bring: at least one deployed tool & three failure cases

SPEAKER NOTES

PREVIEW Set the tone. Week 4 is a different beast: less guided build, more peer review and verification design. You want them excited but realistic about the prerequisite.

Confirm the date, time, and location of Week 4 if you have them. Note that anyone who can't finish Build #3 by then should reach out — there's a path to catch up.

Ship the tool. Update the map.

Questions before we close?

SPEAKER NOTES

CLOSE Open for questions. Don't rush the room out — the questions you get in the last five minutes are often the ones that matter most for what they actually build between now and Week 4.

INSTRUCTOR HOUSEKEEPING — AFTER CLASS Capture the consolidated frontier-map entries from today and push them to the shared doc. Ping any student who didn't ship a working Build #3 to schedule a 1:1 catch-up before Week 4. Note any access issues that came up in Module 1 and report them to IT so the next session starts clean.